

# Operating Systems



Processor Virtualization via Process Abstraction (Lecture-2)

Monsoon 2014, IIIT-H, Suresh Purini

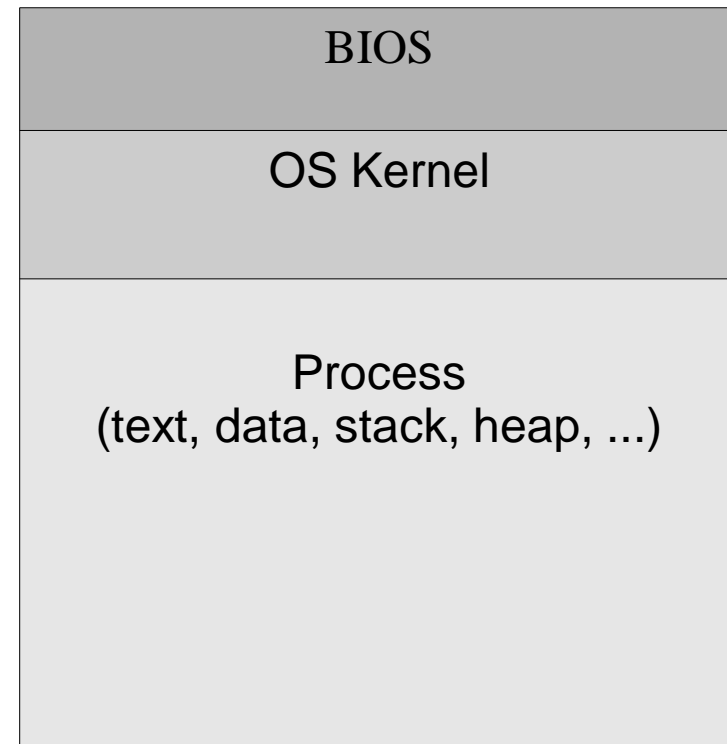
# Virtualizing a Processor

- Goal
  - We would like to run multiple programs on a single CPU (may be more).
  - Each program should get an illusion that it has the entire CPU for itself.
    - In other words, provide each program a separate virtual CPU.
- Possible Tradeoff: Virtual CPU may run a bit slower than the underlying Physical CPU.
- Idea: A program relinquishes the CPU in either of the following two scenarios
  - Has to wait on an I/O request
  - Already used the allocated time slice (time multiplexing)

# Process Abstraction

- A process is a dynamic object which can be created and destroyed using an OS API (aka system calls).
- The state of the process and its dynamics (or evolution) are a function of the static program from which it is created.
- **Simply Put:** A process is a program in execution.
  - A program is a static entity whereas a process is a dynamic entity.
- **Process State:** text/code, data, bss, stack, heap, CPU registers, program counter, ..., anything else!
- As a process executes its state keep evolving.

Memory Map

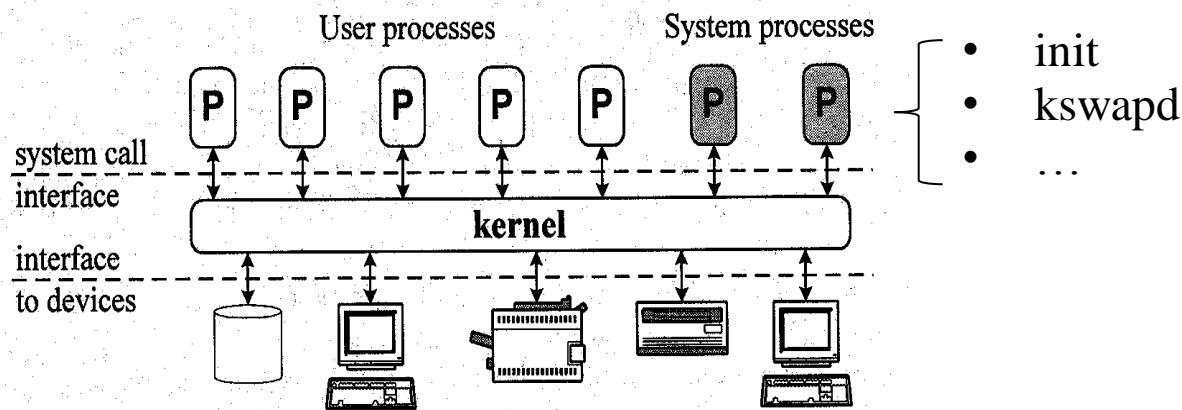


# Per Process Kernel Data Structures


- u-area (accessible only when the process is currently running)
  - Process control block – Stores the hardware context when the process is not running
  - Open file descriptor tables
  - CPU and other resource usage statistics
  - Per kernel stack
  - .....
- proc structure (accessible at any time)
  - Process id
  - Location of the kernel address map for the u-area of this process
  - Current process state (like READY, STOPPED, BLOCKING, ...)
  - .....

# Process Abstraction

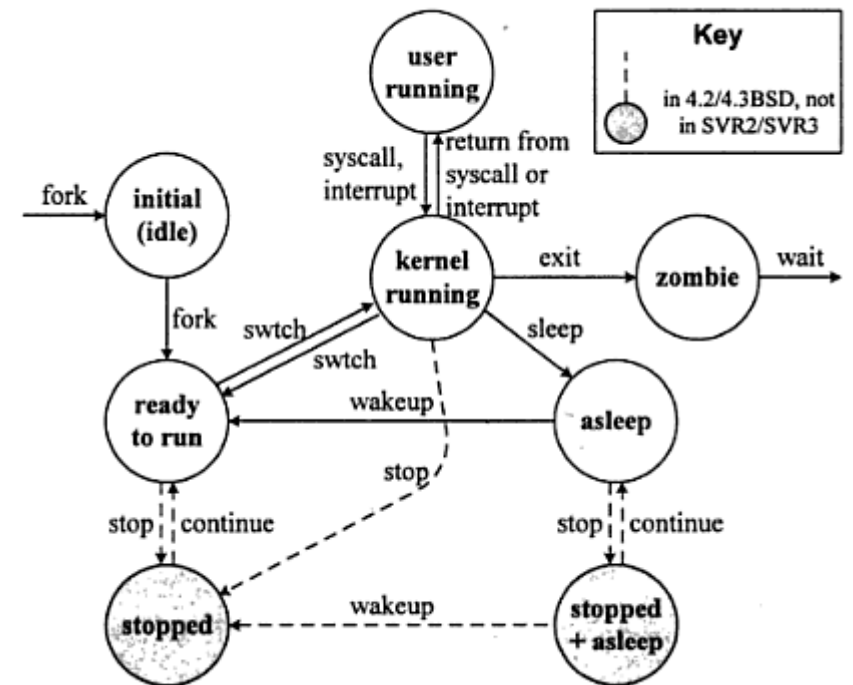

- Every process should get
  - a virtual CPU for itself
  - a complete virtual memory address space
- Different processes may run different instances of the same program
  - E.g., Two Firefox processes are fired from the same program.
- Every process gets a unique identifier (pid)



# Context Switching



# Process State Transitions



**Scheduling Policy:** Among the available ready-to-run processes which one should be allocated CPU in the next time slice.

# Unresolved Issues

1. Kernel isolation
2. Process isolation
3. Address space isolation



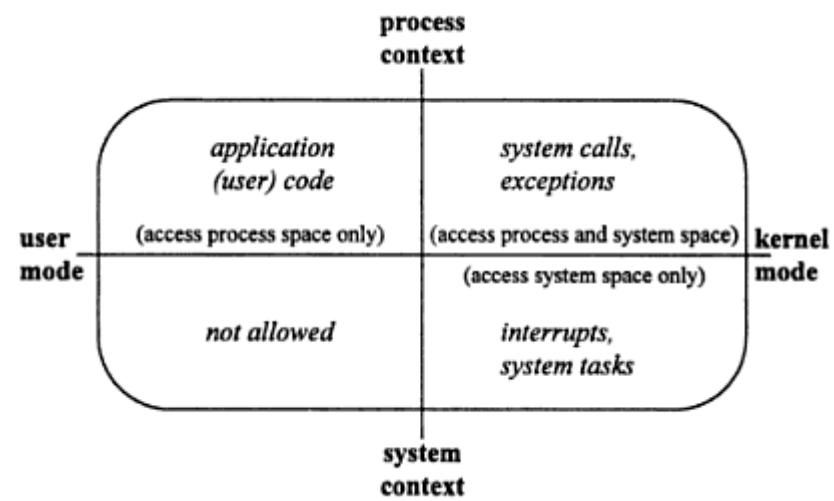
# Tying it All Together: The Unix Shell

```
while(! EOF) {  
  read input  
  handle regular expressions  
  int pid = fork();          // create a child  
  if(pid == 0) {              // child continues here  
    exec("program", argc, argv0, argv1, ...);  
  }  
  else {                      // parent continues here  
    ...  
  }
```

- ◆ Translates <CTRL-C> to the kill() system call with SIGKILL
- ◆ Translates <CTRL-Z> to the kill() system call with SIGSTOP
- ◆ Allows input-output redirections, pipes, and a lot of other stuff that we will see later

# Mode, Space and Context

- User Mode vs Kernel Mode
  - Realized using processor modes
- User Space vs System Space
- Process Context vs System Context



# Preemptive versus Nonpreemptive