

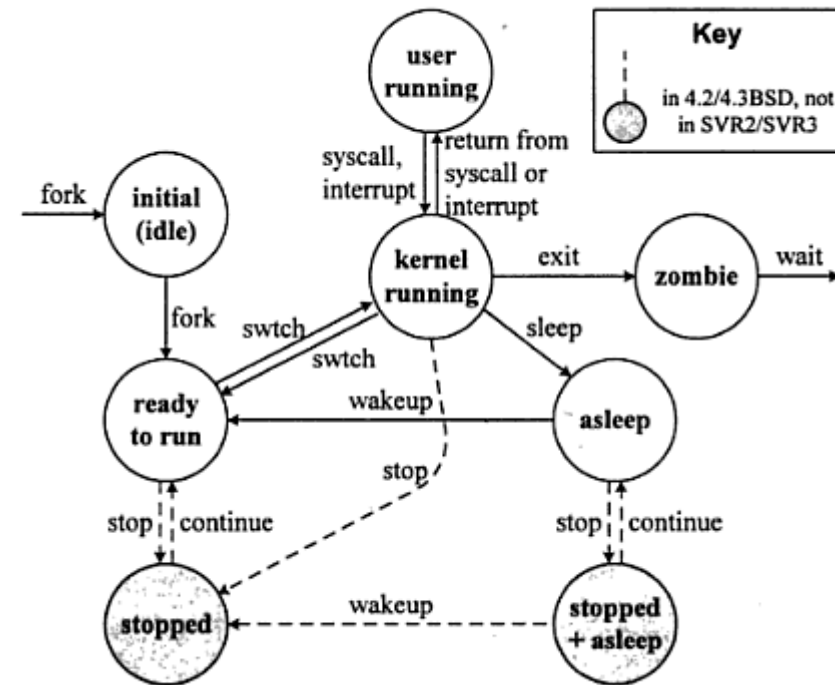
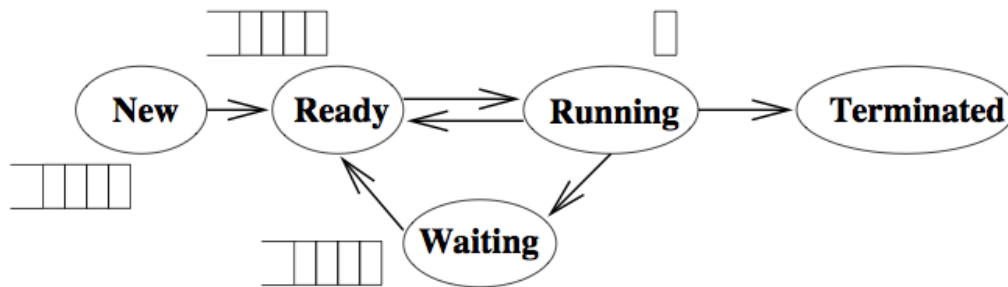
Operating Systems



Scheduling (Lecture-3)

Monsoon 2014, IIIT-H, Suresh Purini

Process States and State Transitions



Two Policy Questions

1. When do we call the scheduler?
2. How does the scheduler chooses the next process?

Mechanism Question: How does the actual context switch happens?

Scheduling the Scheduler

The kernel runs the scheduler whenever

- Some process voluntarily relinquishes the CPU
 - I/O block, termination, SIGSTOP, sleep,
- a processes time slice is over (timer ISR)
-
-
- **Question:** What if a process is in the kernel mode and its time slice is over?

Non-preemptive Kernels

- **Non Pre-emptive Kernels:** A process executing in kernel mode cannot be preempted by another process even though its time quantum expires.
- **Pre-emptive Kernels and Synchronization Issues**

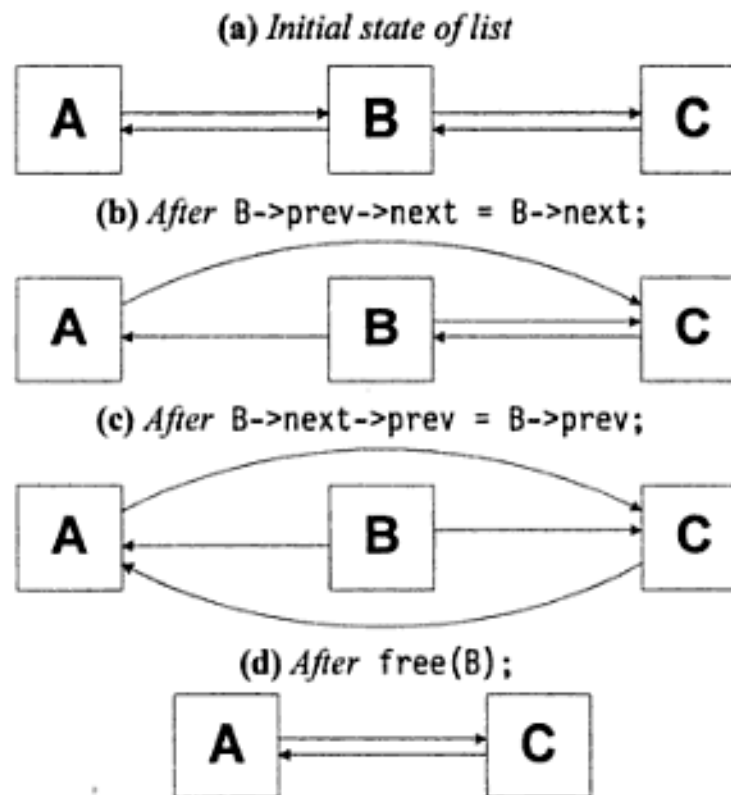


Figure 2-6. Removing an element from a linked list.


Workloads and Scheduling Metrics

- Interactive Applications – Shells, editors, etc.
 - Response time: (first response time – arrival time)
- Batch Applications – scientific computations, kernel compilation, etc.
 - Turnaround time: (completion time – arrival time)
- Real-time Applications – Audio/Video players, missile tracking processes, etc.
 - Strict deadlines

Batch Applications

Assumptions

- All jobs arrive at the same time
- We know the job times ahead
- No I/O blocking (not necessary)
- What should be the **scheduling policy** to minimize the turnaround time?




SJF


Batch Applications

Assumptions

- We know the job times ahead but jobs can have different arrival times
- No I/O blocking (not necessary)
- What should be the **scheduling policy** to minimize the turnaround time?



Convoy Effect





FIFO

Batch Applications

Assumptions

- We know the job times ahead but jobs can have different arrival times
- No I/O blocking (not necessary)
- What should be the **scheduling policy** to minimize the turnaround time?




STCF


Optimizing for Response Time

Assumptions

- We know the job times ahead but jobs can have different arrival times
- No I/O blocking (not necessary)
- What should be the **scheduling policy** to minimize the response time?
 - How would it impact the turnaround time?



SJF

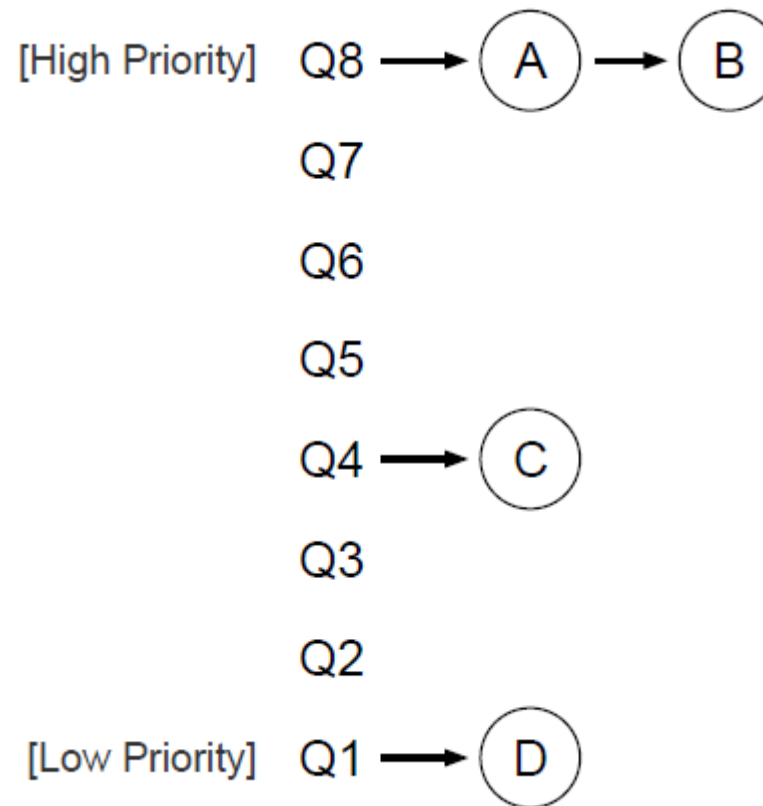


Round Robin

As time slice length increases the amortized cost decreases.

Multi-Level Feedback Queues

- **Rule 1:** If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).
- **Rule 2:** If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR.



- Any problems here?

Multi-level Feedback Queues

- Rule 3: When a job enters the system, it is placed at the highest priority (the topmost queue).
- Rule 4a: If a job uses up an entire time slice while running, its priority is *reduced* (i.e., it moves down one queue).
- Rule 4b: If a job gives up the CPU before the time slice is up, it stays at the *same priority level*.
- Any problems here?




Figure 8.2: Long-running Job Over Time

Mixture of Interactive Jobs and Long Running Jobs


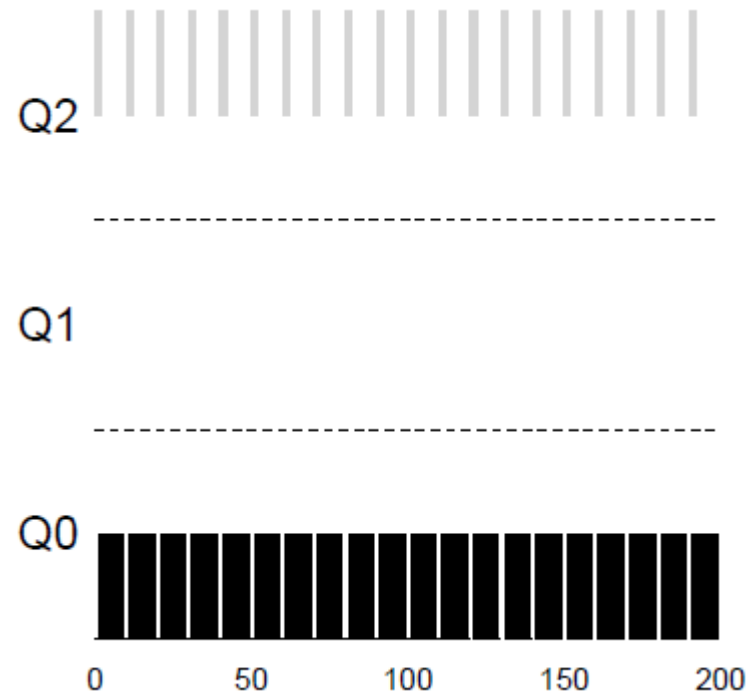



Figure 8.3: Along Came An Interactive Job

Mixture of I/O Intensive and CPU Intensive Workloads




So Far So Good


What is happening here?



Priority Boosting



Gaming the Scheduler



Multi-level Feedback Queues

- **Priority Levels:** A priority level is assigned to every process (let us say between 0 and 100).
 - Lower value means higher priority
 - Processes stuck in the kernel gets high priority (let us say between 0 and 25)
- Always the process at the highest priority gets the CPU.
 - If more than two processes at this level, do round robin.
- Dynamic priority assignment
 - $\text{Process priority} = \text{base priority} + \text{nice value} + (\text{"recent CPU usage"}/\text{constant})$