

Operating Systems

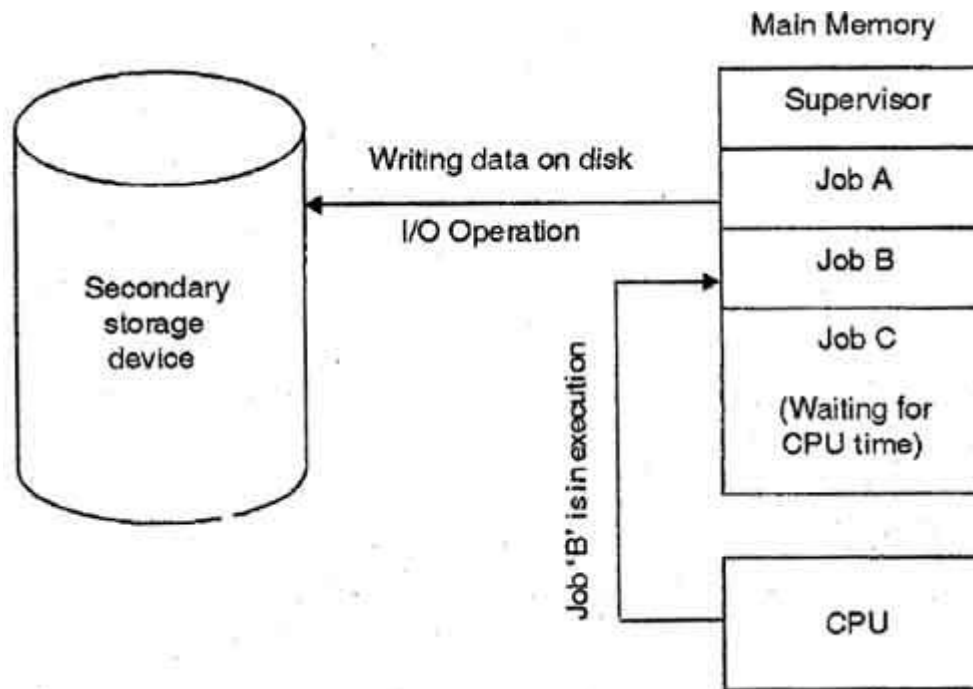


Processor Virtualization via Process Abstraction (Lecture-2)

Monsoon 2014, IIIT-H, Suresh Purini

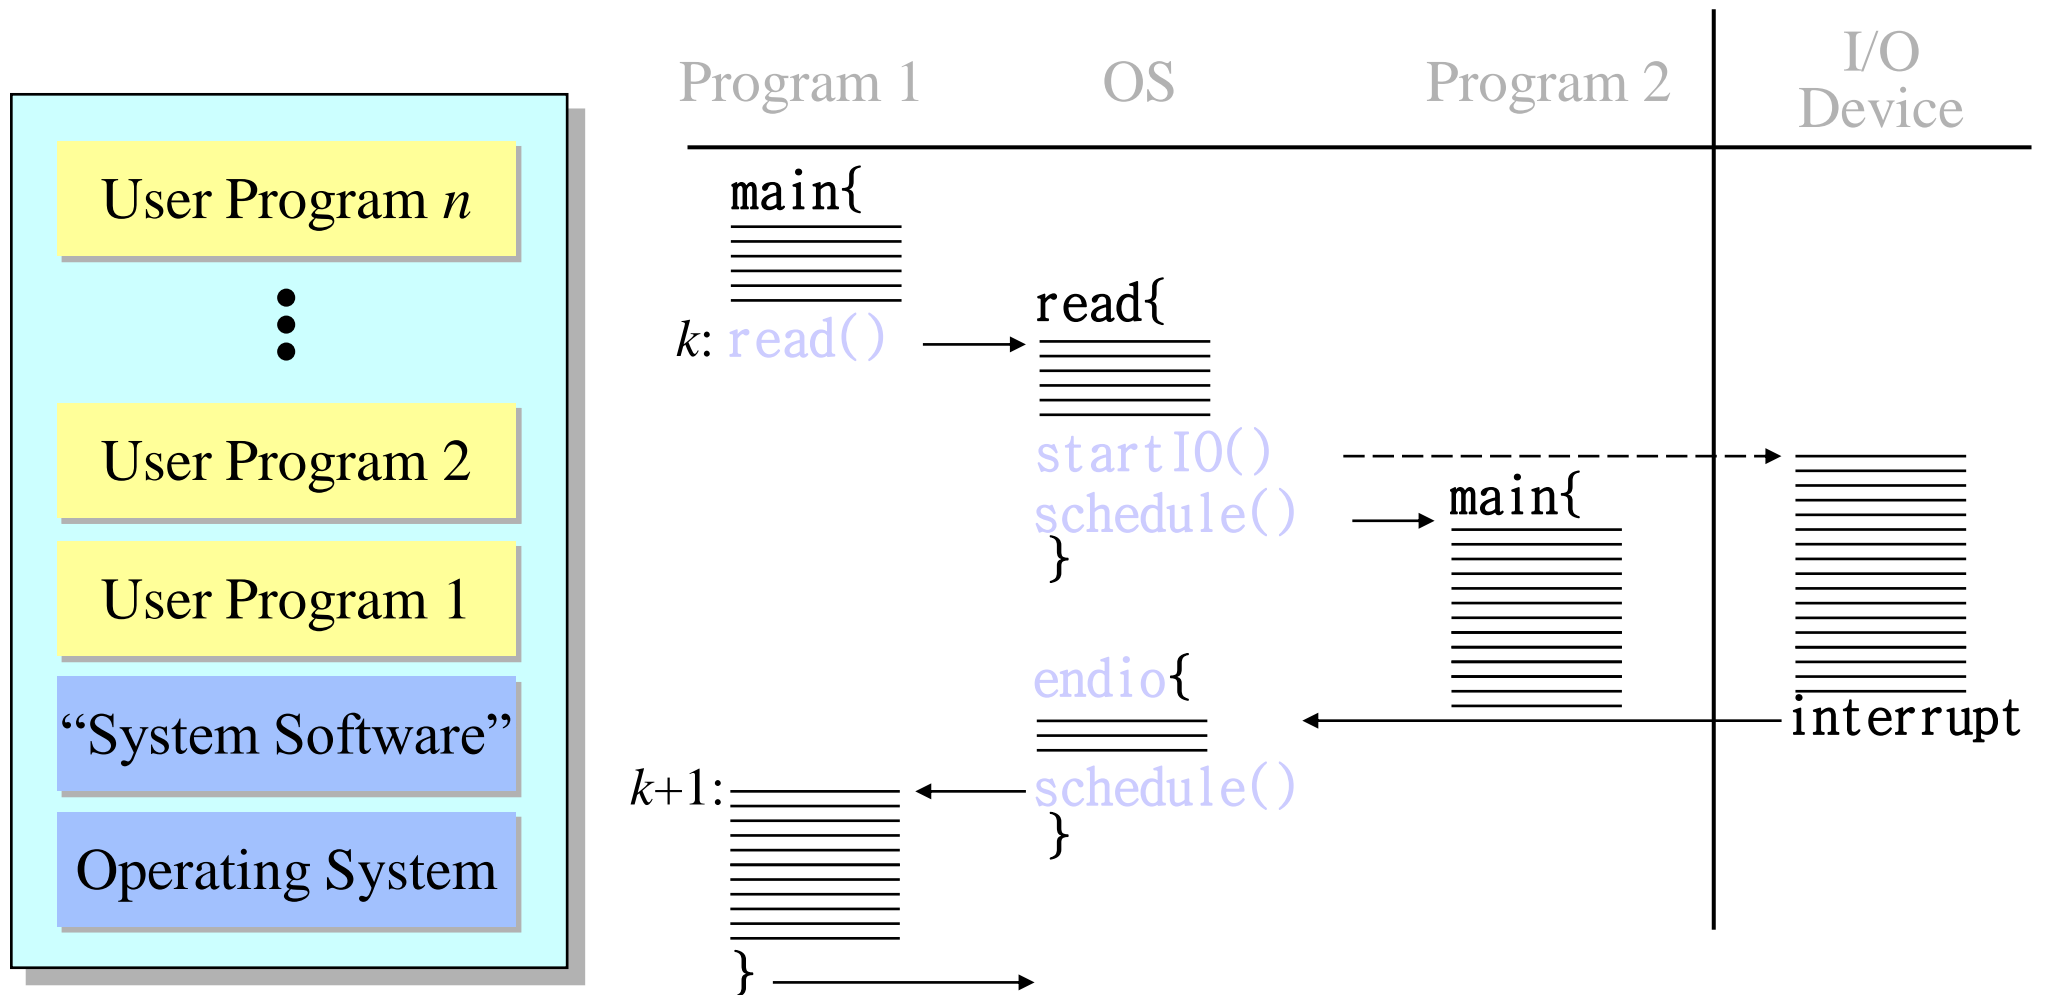
Multiprogrammed Operating Systems

- **Idea:** Switch between multiple jobs based on I/O requests.

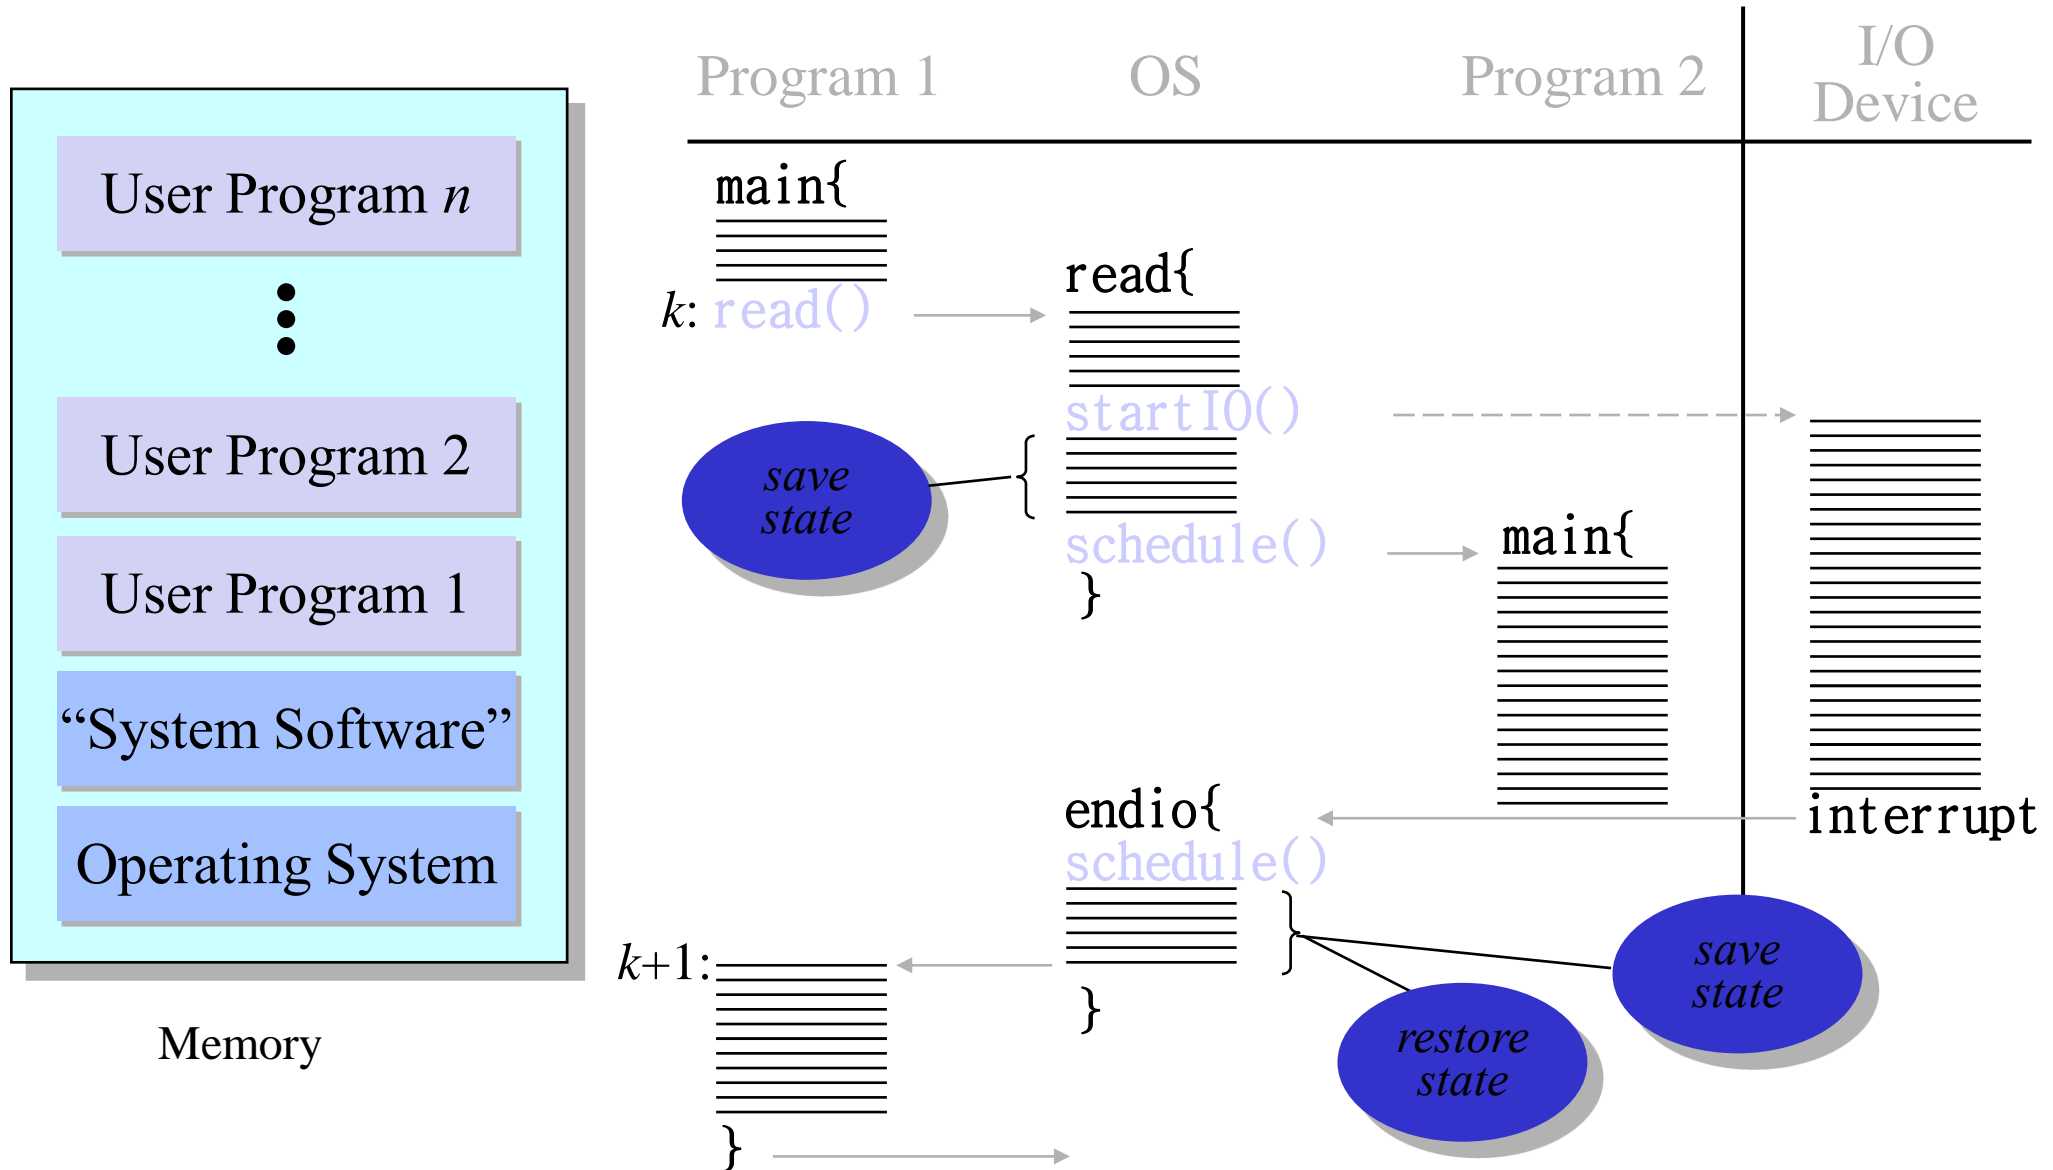


Multiprogramming ('65-'80)

Keep several jobs in memory and multiplex CPU between jobs

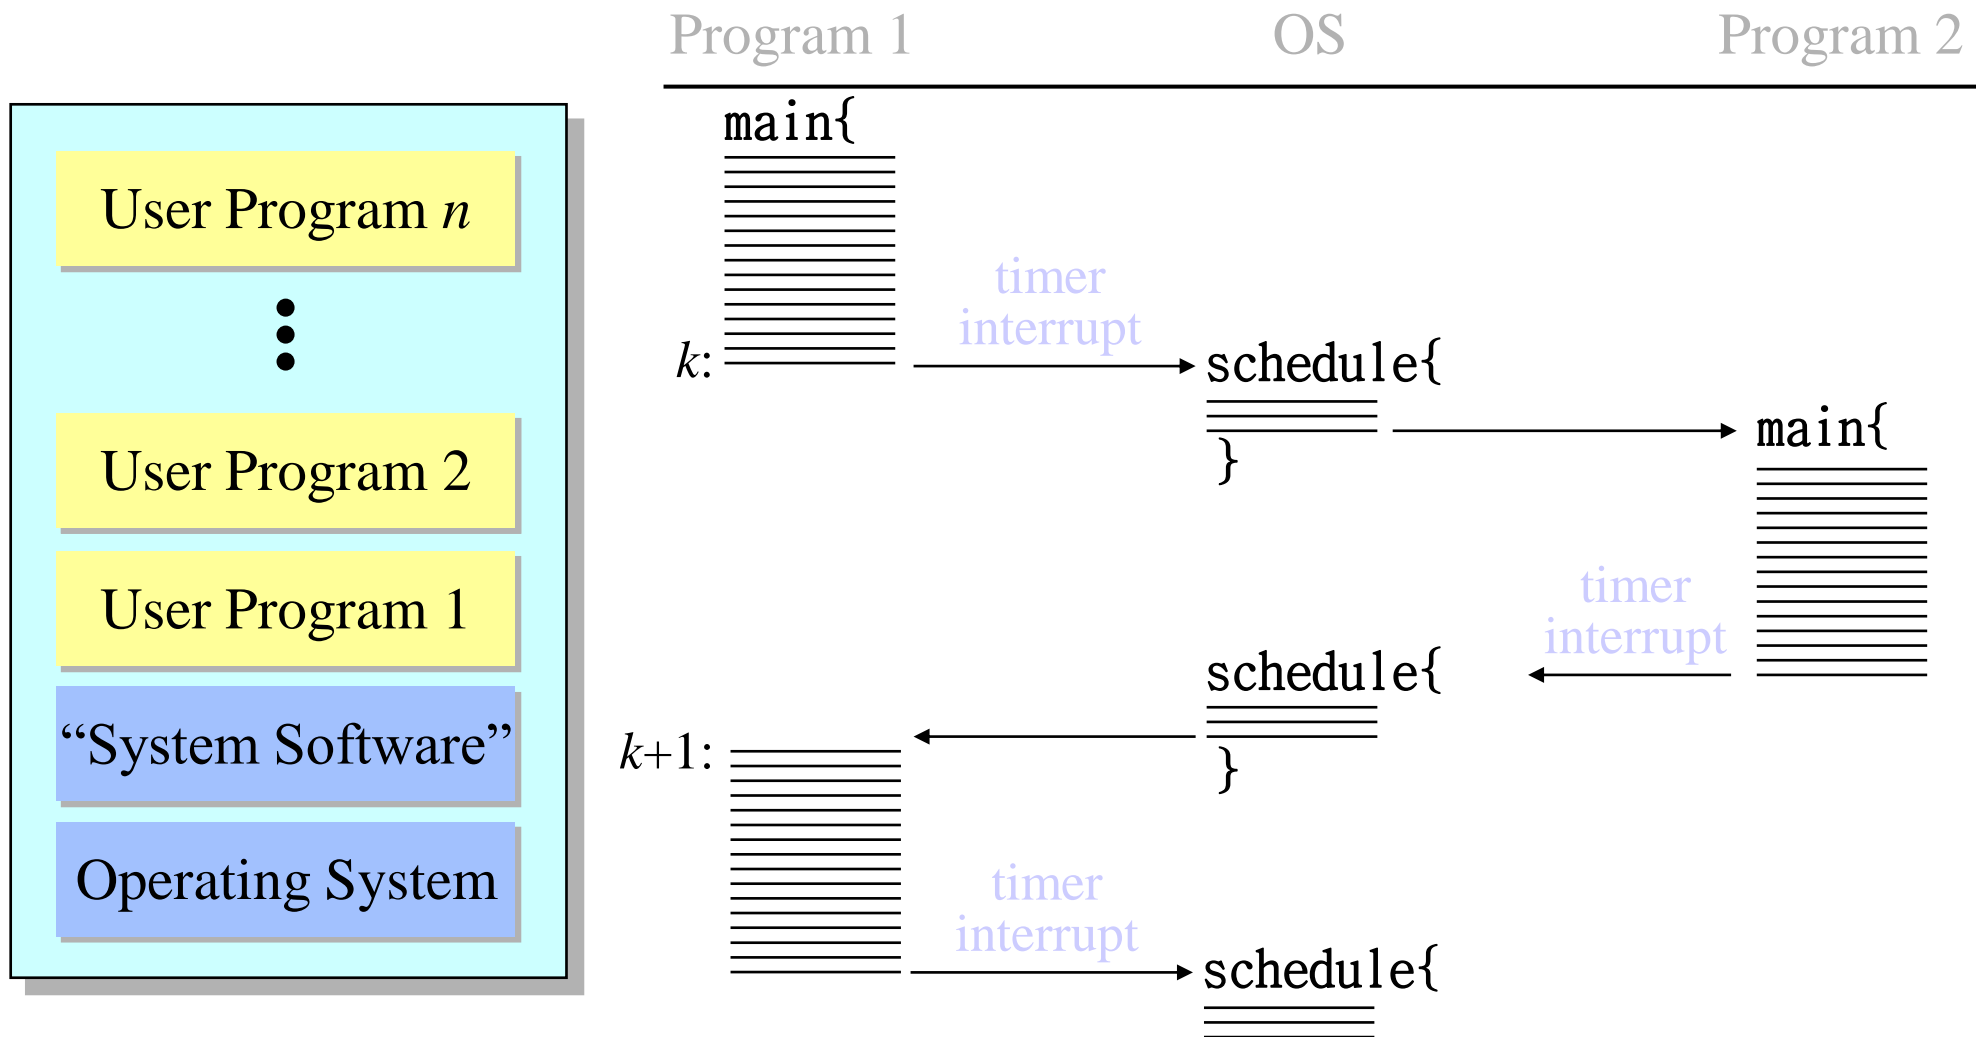


Context Switching

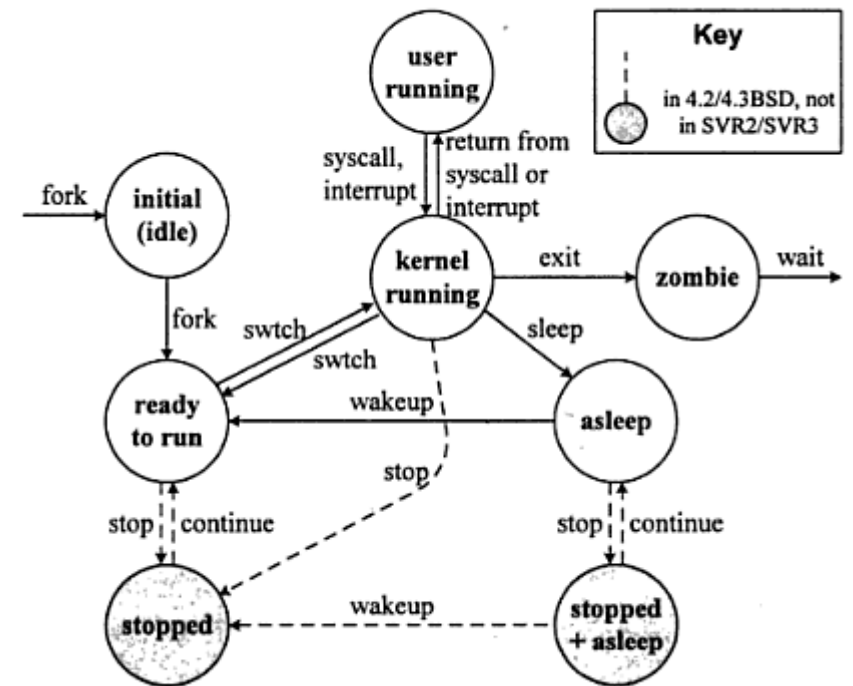
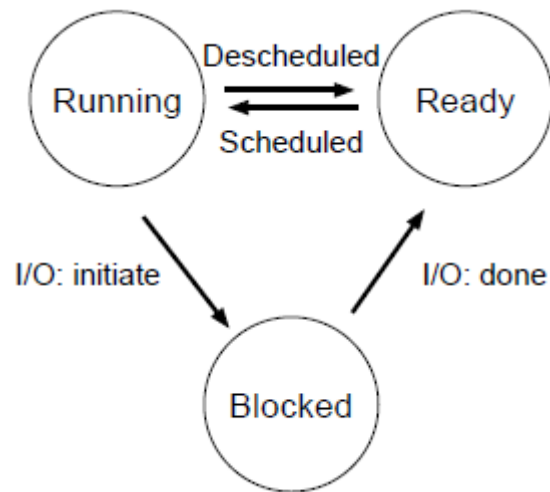


Timesharing Operating Systems ('70-)

A timer interrupt is used to multiplex CPU among jobs



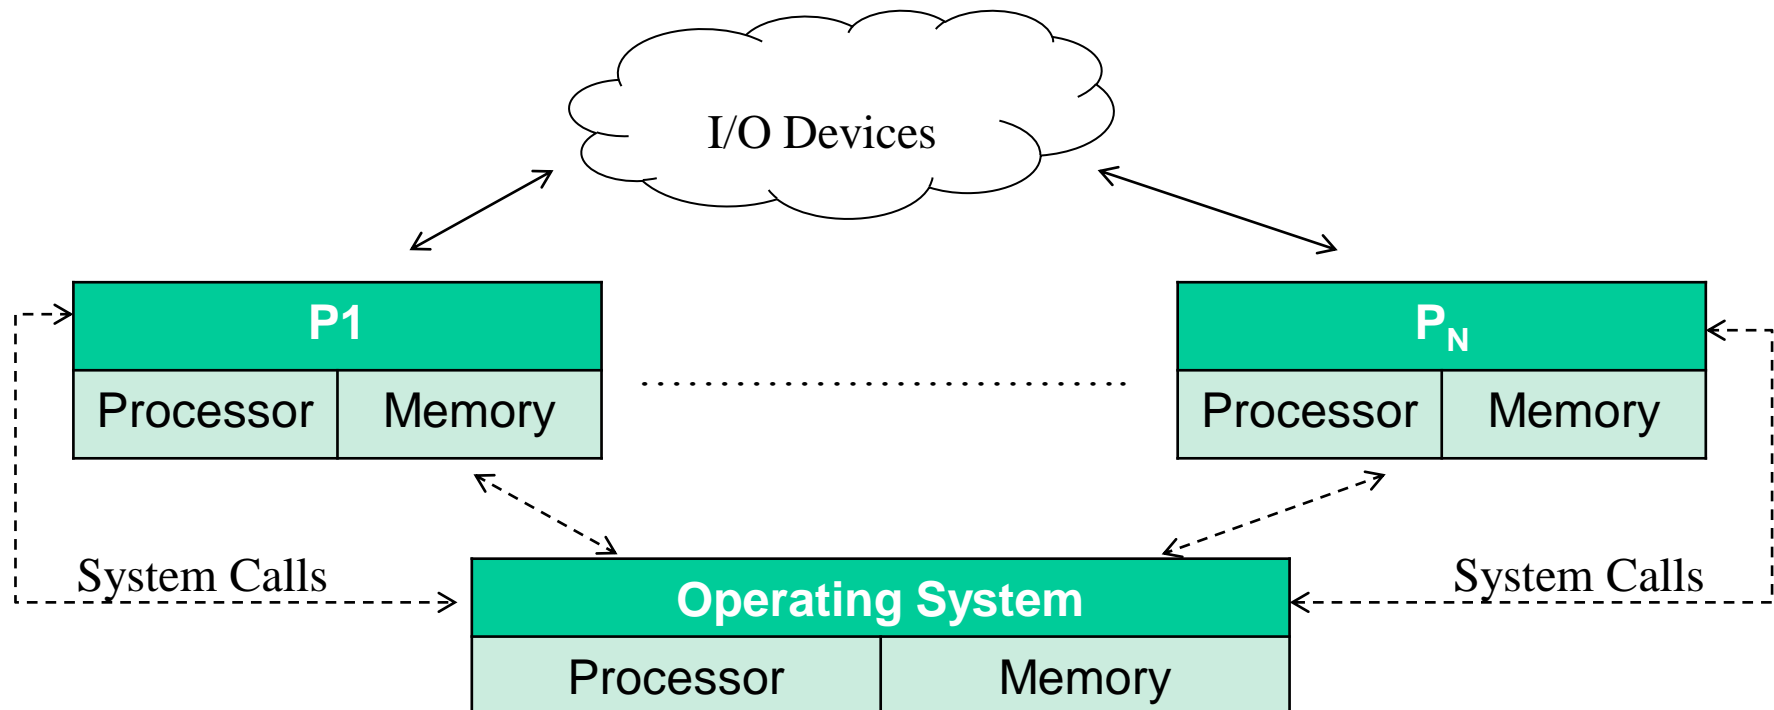
Process State Transitions



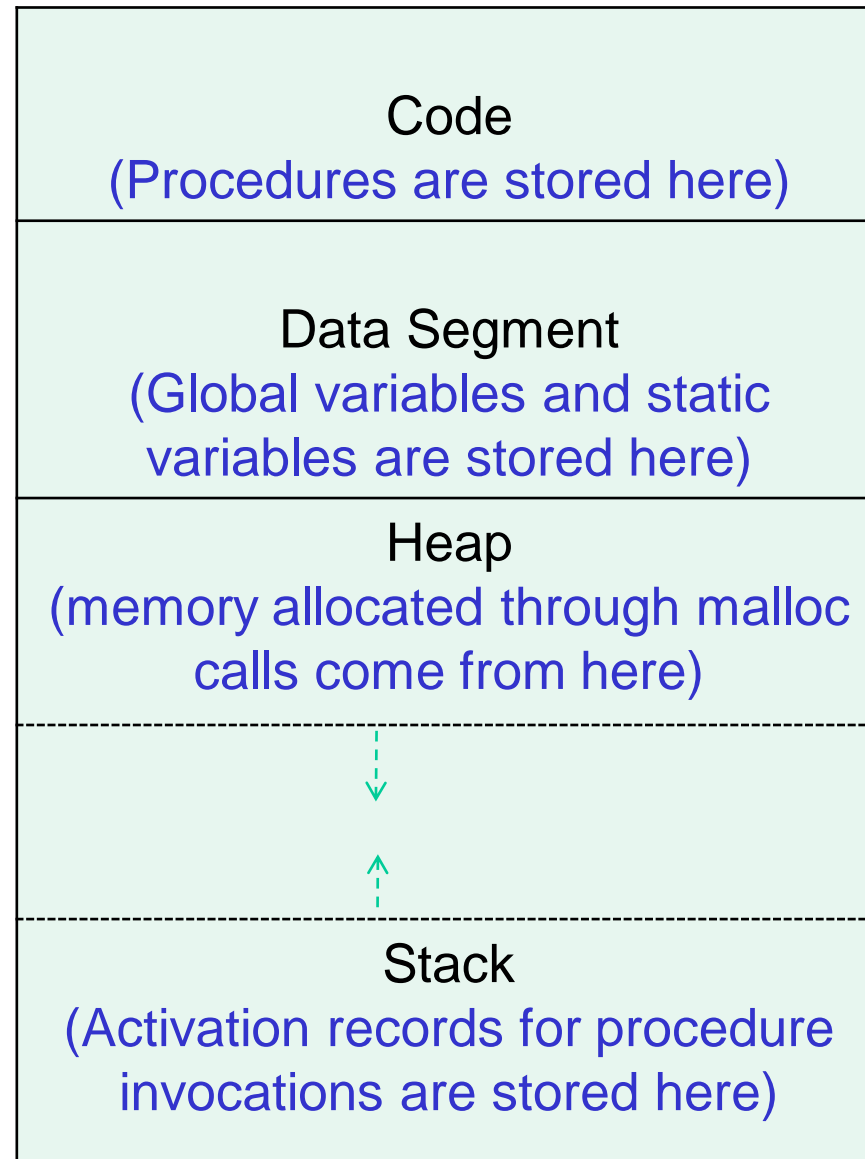
Scheduling Policy: Among the available ready-to-run processes which one should be allocated CPU in the next time slice.

Virtualization 1.0

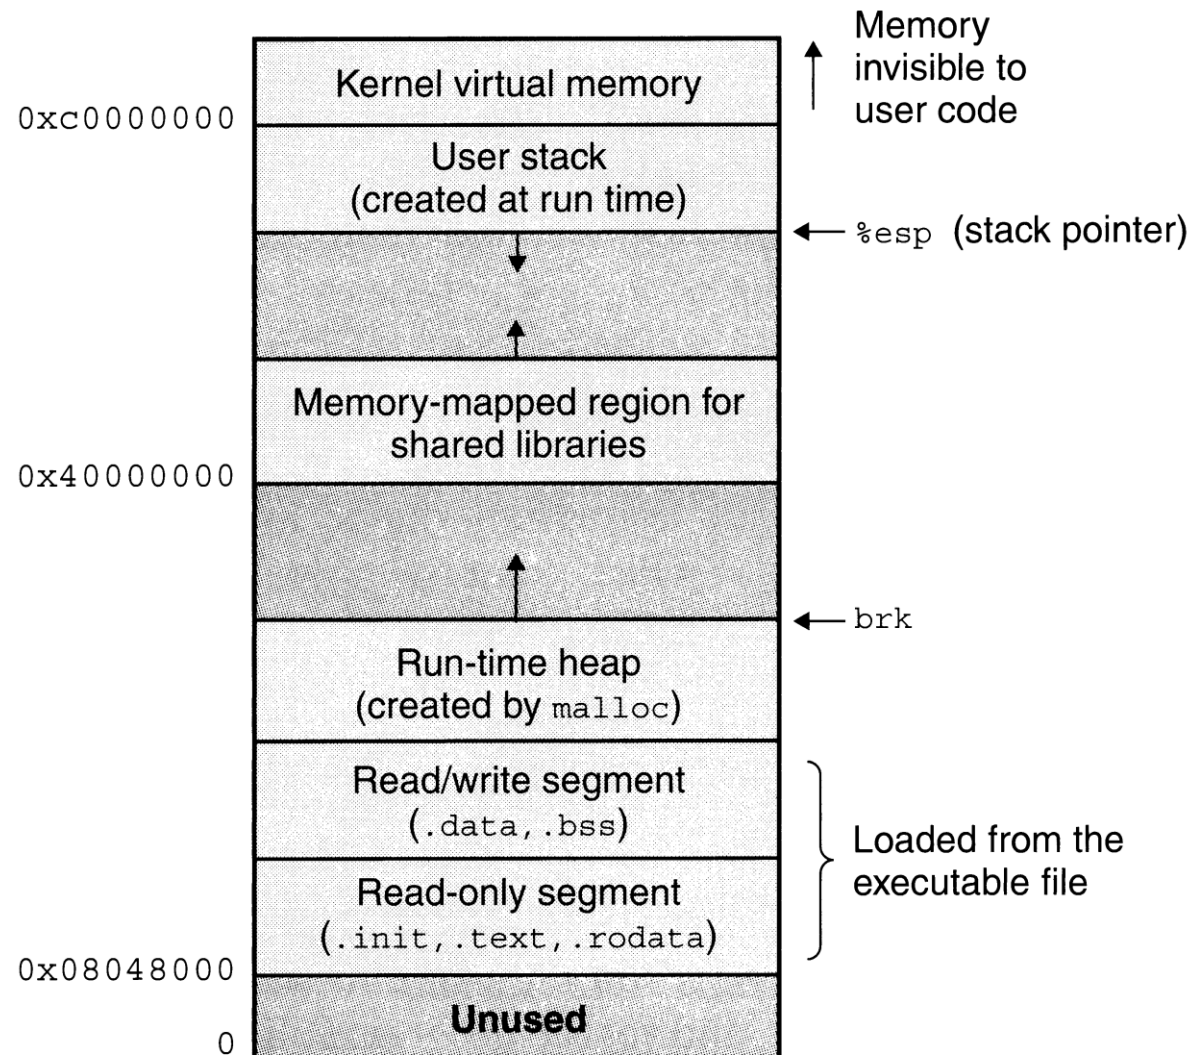
- **Goal:** Give every program a virtual PC
- Every process
 - Gets an illusion that it has an entire processor for itself (via multiplexing CPU over time)
 - Complete memory address space available (via Virtual Memory)
- **However:** Indirect interference between processes is possible



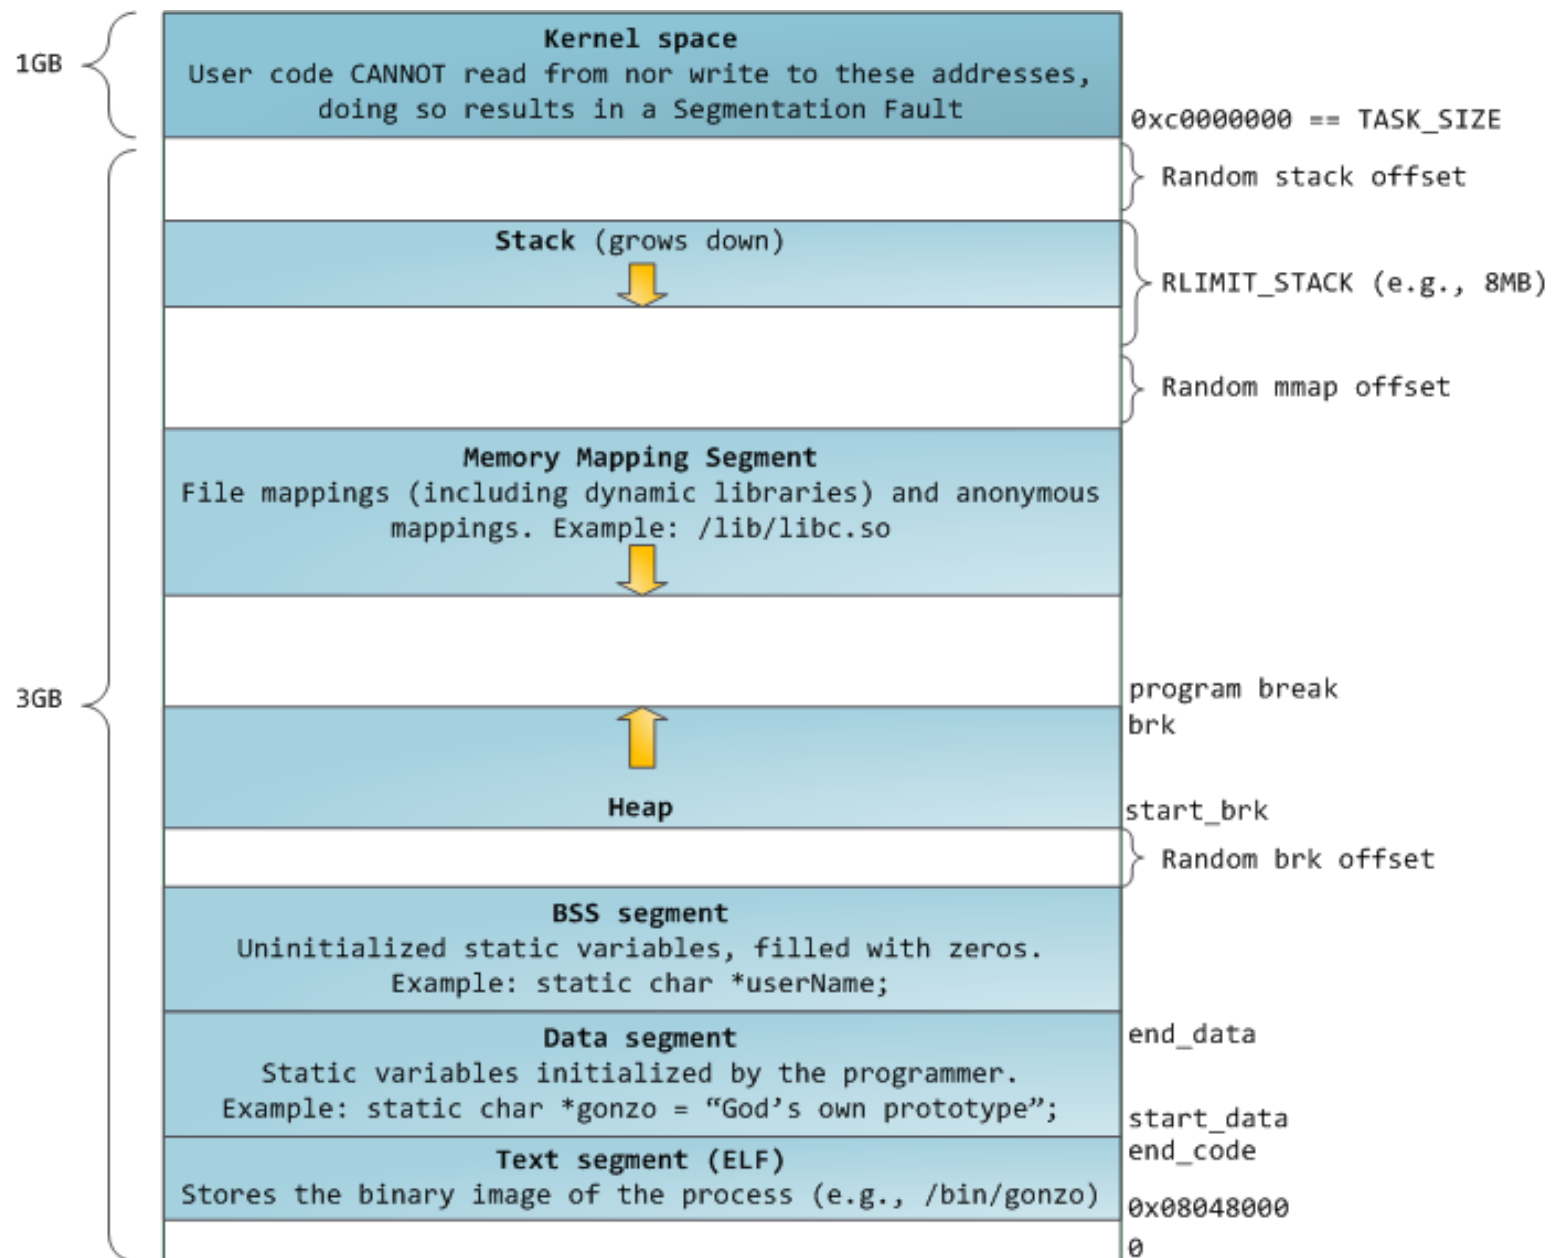
Storage Lay-out of a Program



Virtual Address Space Layout of a Linux Process

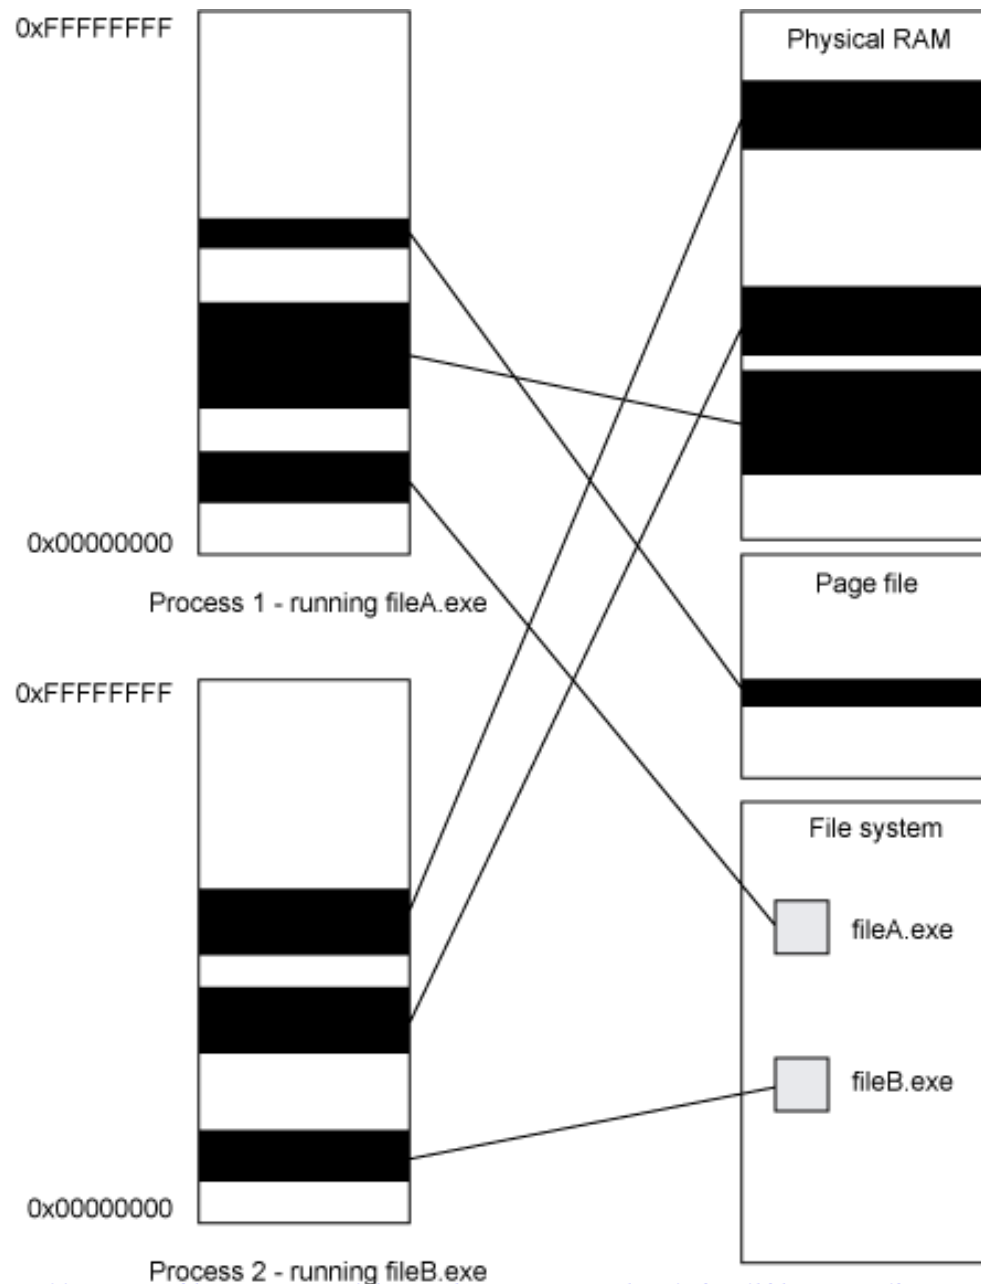


Virtual Address Space Layout of a Linux Process



Taken from: <http://duartes.org/gustavo/blog/post/anatomy-of-a-program-in-memory>

Virtual Address Space versus Physical Address Space



Taken From: <http://www.ibm.com/developerworks/aix/library/j-nativememory-aix/index.html>

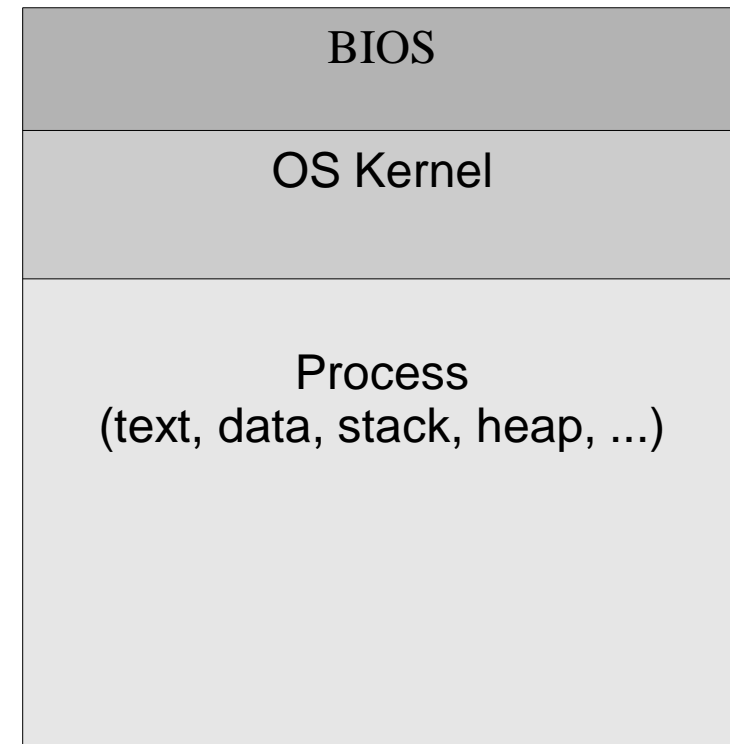
Virtualizing a Processor

- Goal
 - We would like to run multiple programs on a single CPU (may be more).
 - Each program should get an illusion that it has the entire CPU for itself.
 - In other words, provide each program a separate virtual CPU.
- Possible Tradeoff: Virtual CPU may run a bit slower than the underlying Physical CPU.
- Idea: A program relinquishes the CPU in either of the following two scenarios
 - Has to wait on an I/O request
 - Already used the allocated time slice (time multiplexing)

Process Abstraction

- A process is a dynamic object which can be created and destroyed using an OS API (aka system calls).
- The state of the process and its dynamics (or evolution) are a function of the static program from which it is created.
- **Simply Put:** A process is a program in execution.
 - A program is a static entity whereas a process is a dynamic entity.
- **Process State:** text/code, data, bss, stack, heap, CPU registers, program counter, ..., anything else!
- As a process executes its state keep evolving.

Memory Map




Per Process Kernel Data Structures

- u-area (accessible only when the process is currently running)
 - Process control block – Stores the hardware context when the process is not running
 - Open file descriptor tables
 - CPU and other resource usage statistics
 - Per kernel stack
 -
- proc structure (accessible at any time)
 - Process id
 - Location of the kernel address map for the u-area of this process
 - Current process state (like READY, STOPPED, BLOCKING, ...)
 -

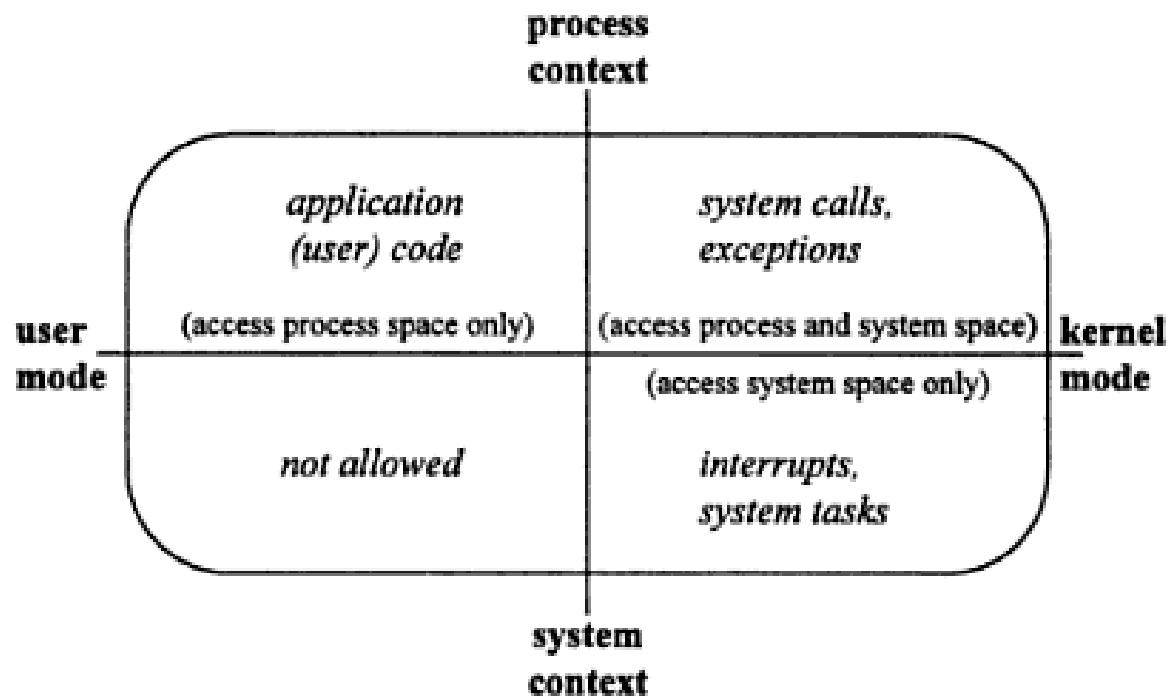
Process Abstraction

- Every process should get
 - a virtual CPU for itself
 - a complete virtual memory address space
- Different processes may run different instances of the same program
 - E.g., Two Firefox processes are fired from the same program.
- Every process gets a unique identifier (pid)



Mode, Space and Context

- User Mode vs Kernel Mode
 - Realized using processor modes
- User Space vs System Space
- Process Context vs System Context



Unresolved Issues

1. Kernel isolation
2. Process isolation
3. Address space isolation

Process Creation

Parent

```
main()    pid = 3456
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}


void ParentProcess()
{
    .....
}
```

Child

```
main()    pid = 0
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```



Process Creation

Parent

```
main()
{
    pid = 3456
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

Child

```
main()
{
    pid = 0
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

Process Creation

Parent

```
main()    pid = 3456
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```


Child

```
main()    pid = 0
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

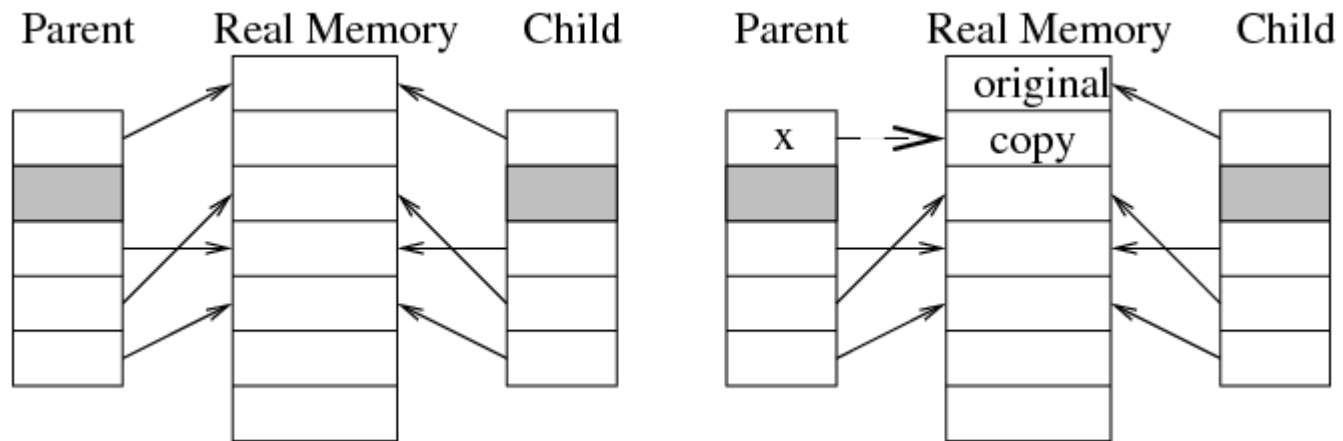
void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```

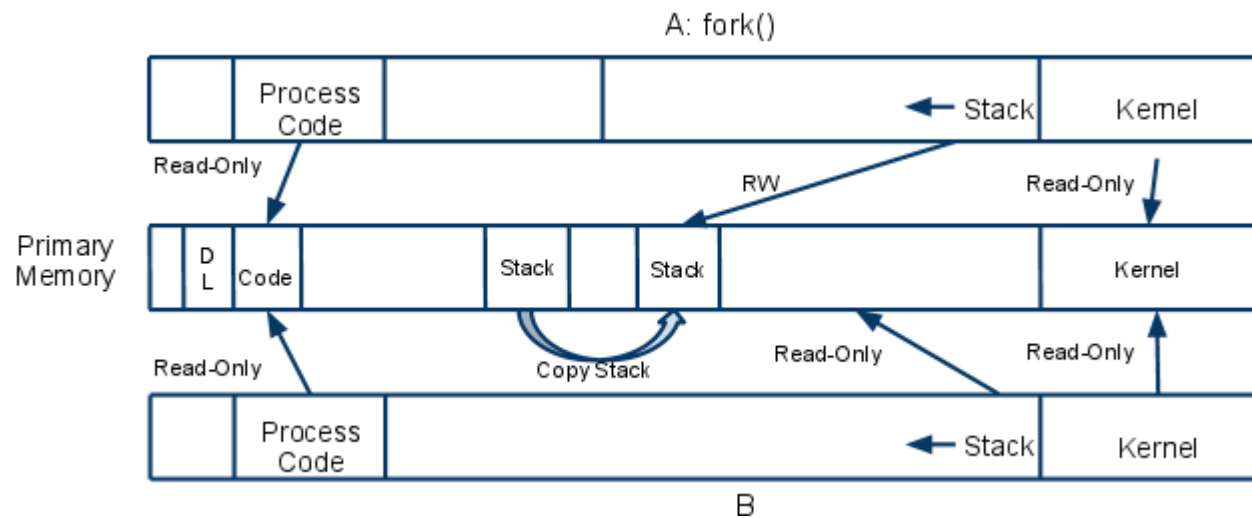
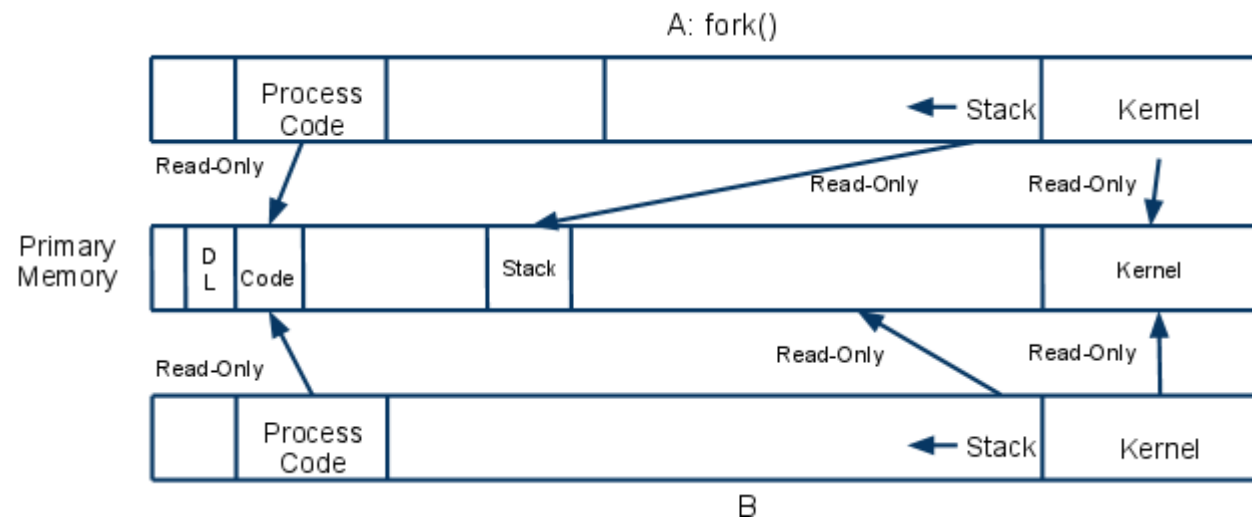
Process Creation



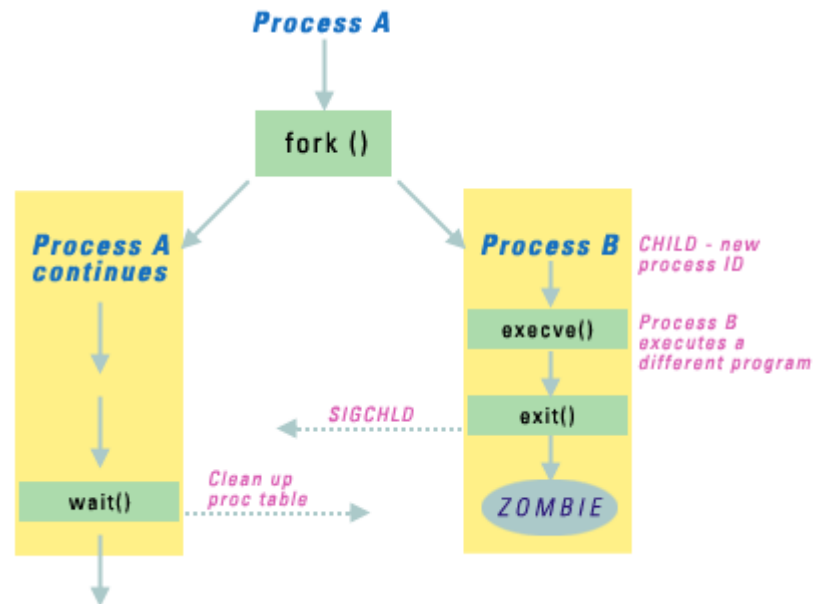
Copy-on-Write






Copy-on-Write



execve



Copy-on-Write



Tying it All Together: The Unix Shell

```
while(! EOF) {  
  read input  
  handle regular expressions  
  int pid = fork();          // create a child  
  if(pid == 0) {              // child continues here  
    exec("program", argc, argv0, argv1, ...);  
  }  
  else {                      // parent continues here  
    ...  
  }
```

- ◆ Translates <CTRL-C> to the kill() system call with SIGKILL
- ◆ Translates <CTRL-Z> to the kill() system call with SIGSTOP
- ◆ Allows input-output redirections, pipes, and a lot of other stuff that we will see later